

K-Means Algorithm with CUDA

Analysis and solutions for dividing an algorithm's workload with the objective of reducing its execution time.

1st Vicente Moreira
Escola de Engenharia
Universidade do Minho
Braga, Portugal
pg50799@alunos.uminho.pt

2nd Joana Alves
Escola de Engenharia
Universidade do Minho
Braga, Portugal
pg50457@alunos.uminho.pt

Abstract—This is the final document in the evolution of the basic clustering algorithm K-means. Here, further possible optimizations, such as the use of a *MPI*, or *CUDA*, will be analyzed along with its implementations and testing. The results will also be compared to the last algorithm iteration developed in the previous document to achieve a final evaluation of the overall project.

Index Terms—K-Means, MPI, CUDA, Blocks, Threads

I. INTRODUCTION

A. Context

In the last part of this project, our team developed a solution where the heavy workload of the k-means algorithm was split between threads among a CPU, reducing its overall execution time. To further improve this program, our team aims to use more efficient parallelization techniques, either by increasing the amount of threads available, without reducing the overhead, or splitting the algorithm tasks between processes.

For the latter the use of a *MPI* (Message Passing Interface) is essential as it allows the communication and transfer of data between processes. With the use of many independent processes, where which one is assigned a specific task inside the algorithm, a *pipeline* can be achieved inside an application, reducing its execution time.

As for the execution of many threads, it's well documented that GPU's are optimized and developed around the computing of large unsegmented blocks of data, having an extreme affinity to parallelization, where a larger number of threads is used to process these blocks of data at the same time with little to no overheads. To access these benefits, an API such as *CUDA* (Compute Unified Device Architecture) will help our team develop a better algorithm.

II. TESTS AND METRICS TO BE EVALUATED

In this part of the project, our team will evaluate the algorithm performance by two main metrics: execution time, and Level 1 Cache Misses, since these are well linked with the overall increase in the performance of an application.

To test the algorithm and its new versions, we will use a new number of Points and Clusters as a baseline between all tests, executing the algorithm using 1 Million Points and 320 Clusters. This change happened due to a restriction present on the use of the *CUDA* version, that will be explained in its

corresponding chapter. We will also measure the application performance for up to 1024 Clusters, to evaluate the solution's scalability.

The tests will be performed using two machines: the SEARCH Cluster, provided by Uminho and used throughout the previous parts of this project and a personal Laptop with the following specs:

Laptop ASUS	
OS	Ubuntu LTS 20.04
CPU	Intel® Core™ i5-8300H CPU @ 2.30GHz × 8
RAM	8GB
GPU	Nvidia GeForce GTX 1050 Mobile

For each set of conditions, our team will execute 3 tests and calculate the average of the metrics, to further reduce variances inside the same machine.

III. INITIAL ANALYSIS

A. Computationally Heavy Functions

Samples: 53K of event 'cycles', Event count (approx.): 3028987

Overhead	Command	Shared Object	Symbol
92,81%	k_means	k_means	[.] atribuiClusters
3,53%	k_means	k_means	[.] ajustaCentroides
1,16%	k_means	libc-2.31.so	[.] __vfscanf_internal
0,64%	k_means	k_means	[.] inicializa
0,10%	k_means	[unknown]	[k] 0xffffffffffb0204e83


By analyzing the computationally heavy functions of our previous sequential version, we know that **atribuiClusters** is the main function of this algorithm since it always has a $N \times K$ calculations per iteration.

By paralelizing this function, using OMP (Open Multi-Processing) we successfully reduced the algorithm execution time but it still remained our applications heaviest function, implying that the only way to improve its execution time in a major way would come from a better parallelization method.

Samples: 62K of event 'cycles', Event count (approx.): 52458939169

Overhead	Command	Shared Object	Symbol
96,22%	k_means	k_means	[.] atribuiClusters_omp_fn.0
2,25%	k_means	k_means	[.] ajustaCentroides
0,60%	k_means	libc-2.31.so	[.] __vfscanf_internal
0,32%	k_means	k_means	[.] inicializa

This is further proven by our previous work, where we also implemented parallelization on the function **ajustaCentroides**, and achieved a smaller performance gain.



B. MPI Struggles

With this analysis in mind, we decided not to develop any version of our algorithm that utilized MPI, since its main strong suit is dividing an algorithm tasks between processes and, by processing the data through "blocks", it allows the creation of an efficient pipeline. In this case, our pipeline would be severely bottlenecked by the function **atribuiCentroids**, resulting in a large amount of wasted time by other processes, and achieving a similar execution time to its sequential version.

Another consideration was utilizing MPI to divide this function workload, and although this is possible, our team concluded that this implementation of MPI would be no different to our parallelized OMP version and, by possibly using larger overheads between processes, our execution time would either match the OMP version or increase it.

C. CUDA

Since our analysis indicates that the only way to significantly increase our performance gain in our algorithm is by utilizing a more efficient parallelization strategy/method, we turned our focus on the use of CUDA, an API capable of using the GPU's highly optimized parallelization data processing.

One restriction that comes with the use of this API, is the use and definition of the number of blocks and number of threads per block that CUDA will use for the execution of a given function. These have a max value of 1024, which means we have a maximum number of 1.048.576 threads. This value later restricted our testing by only allowing our team to use up to 1M points in the algorithm.

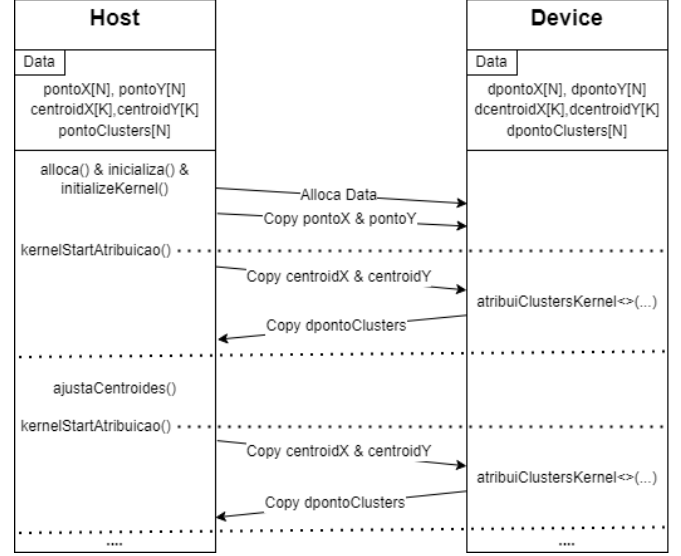
IV. FIRST VERSION DEVELOPMENT

A. Strategy

For our first version, we started by applying CUDA only on the **atribuiClusters** function. For this we needed to start by allocating all data arrays related to the points, centroids and cluster attribution on the device since they would all be needed for this phase. We also use this initialization process to copy the points data to the device, since this data is never changed in-between iterations.

After that, for each time the cluster attribution phase is called, we update the centroids data to the device, since the coordinates may have changed and, after executing the kernel function, we copy the results back to the host.

B. Architecture



C. Results

CudaV1		
Test	TExec (s)	L1 Cache Misses
N=1M, K=320, @SEARCH	1,57	1,19E+07
N=1M, K=1024, @SEARCH	1,78	2,29E+07
N=1M, K=1024, @Laptop	0,81	4,42E+07

(Further result analysis and comparison between versions will be discussed in the corresponding chapter)

D. Profiling

With the use of **nvprof** (Nvidia Profiler) and **nvvp** (Nvidia Visual Profiler), we analysed our profiling metrics to see how the API calls to CUDA were handled and how much overhead might have been introduced by this method.

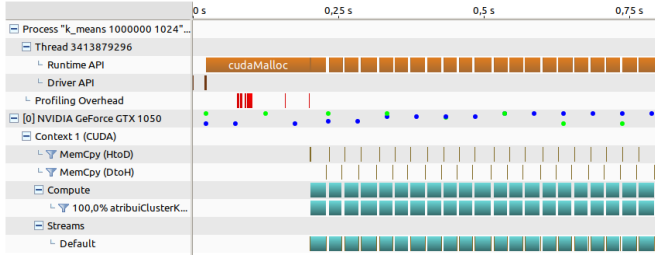
Due to the missing GUI in the SEARCH Cluster, needed for the use of **nvvp**, we will be showing the profilings from the Laptop test:

```

==27541== NVPROF is profiling process 27541, command: ./bin/k_means 1000000 1024
==27541== Profiling application: ./bin/k_means 1000000 1024
==27541== Profiling result:

```

Type	Time(%)	Time	Calls	Avg	Min	Max	Name
GPU activities:	96.71%	486.67ms	21	23.175ms	21.604ms	26.688ms	atribuiClustersKernel
	3.02%	15.197ms	21	723.65us	686.33us	806.37us	[CUDA memcpy DtoH]
	0.27%	1.3530ms	44	30.770us	1.2160us	606.65us	[CUDA memcpy HtoD]
API calls:	74.26%	508.11ms	65	7.8170ms	7.1730us	27.565ms	cudaMemcpy
	25.55%	174.97ms	5	34.995ms	4.1610us	174.58ms	cudaMalloc
	0.16%	694.45us	5	138.89us	6.5130us	202.10us	cudaFree
	0.08%	540.02us	21	25.715us	22.131us	44.947us	cudaLaunchKernel
	0.03%	227.02us	97	2.3400us	170ns	98.253us	cuDeviceGetAttribute
	0.02%	153.26us	1	153.26us	153.26us	153.26us	cuDeviceGetAttribute
	0.01%	40.216us	1	40.216us	40.216us	40.216us	cuDeviceGetName
	0.00%	23.545us	66	356ns	172ns	900ns	cudaGetLastError
	0.00%	8.2200us	1	8.2200us	8.2200us	8.2200us	cuDeviceGetPCIBusId
	0.00%	2.3840us	3	794ns	185ns	1.9380us	cuDeviceGetCount
	0.00%	934ns	2	467ns	188ns	746ns	cuDeviceGet
	0.00%	318ns	1	318ns	318ns	318ns	cuDeviceGetUuid



(Due to the way the profilers work, both profilings correspond to different executions so their data may not match)

From this analysis we can spot a few relevant points:

- The **cudaMalloc** API call has a small but noticeable overhead, taking around 177ms to process.
- We can see there is 21 calls to the function **atribuiClustersKernel**, which matches the first execution plus the 20 pre-defined iterations. These in total take around 500ms to execute, averaging 24ms per call of this function.
- The two data copies combined take around 16ms to execute. 750us per iteration.
- By using the RuntimeAPI time in **nvvp** (800ms) we can subtract the **cudaMalloc** and **atribuiClustersKernel** execution time and obtain the overhead/downtime that our application has due to its **ajustaCentroids** function executing on the Host.
 $800\text{ms} - 500\text{ms} - 177\text{ms} = 123\text{ms of CUDA downtime.}$
 $123\text{ms} - 16\text{ms} = 107\text{ms of Host processing Time.}$
- We can confirm (via **nvvp**) that before the kernel execution, there is a 4096kB data copy (HtoD) that occurs twice, corresponding to the centroid data.
 $(K * \text{sizeof(float)}: 1024 * 4 = 4096\text{kB})$
- We can confirm (via **nvvp**) that after the kernel execution, there is 4MB data copy (DtoH), corresponding to the cluster attribution data.
 $(N * \text{sizeof(int)}: 1.000.000 * 4 = 4\text{MB})$

V. SECOND VERSION DEVELOPMENT

A. Strategy

For the next iteration, we aim to reduce the need to transfer data between host and device in every iteration. This data transfer is needed since the host needs the last cluster attribution to execute the **ajustaCentroids** function and, after adjusting the centroids, this data needs to be transferred to the device again to execute the next attribution.

To avoid this, and to increase the applications parallelization, we want to execute the **ajustaCentroids** function utilizing the CUDA capabilities. In this function, we use three auxiliary arrays with K size, two to measure the sum of the X and Y coordinates and 1 to keep track of how many points the cluster has. To adapt this function into it's CUDA implementation we will need these 3 arrays on memory, increasing the data usage on the Device. (12 * K bytes)

One of the challenges introduced in the CUDA implementation is the number of Threads to use in this section. If we

use K number of Threads, all of them we will have to iterate over every point, which is slower and undesired. If we use N number of Threads, we need to utilize a synchronizing method, so we can "reset" all the auxiliary arrays, before starting the sums and, before calculating the new centroid coordinate, we have to wait for all the threads to finish the sum process. This implementation will also need some data race control process, since some threads might write to the same array location.

B. Synchronization

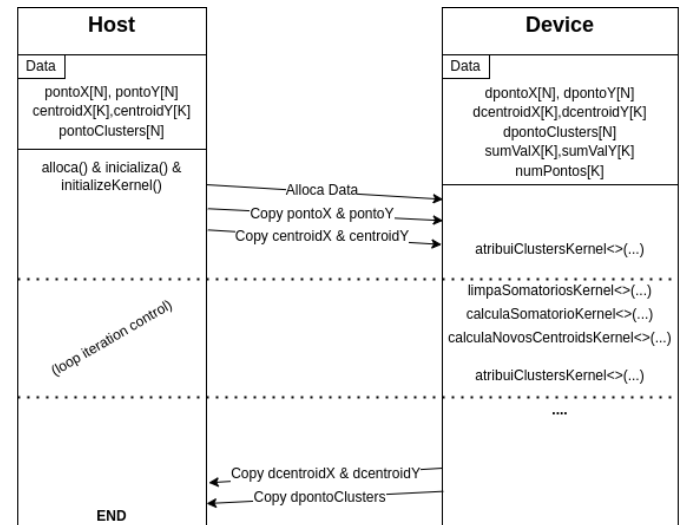
Our team planned to use the CUDA call '**__syncthreads()**' to handle the synchronization process but after some research on the documentation provided, we found out that this function only synchronizes threads within a block, which doesn't work for what we intend, since we need every thread to wait for the "resetting" of the auxiliary arrays (this problem also happens with the final centroid coordinate calculation phase).

The solution we found is commonly named "CPU synchronization", where we divide the **ajustaCentroid** function in many smaller Kernel functions. (**limpaSomatoriosKernel**, **calculaSomatorioKernel**, **calculaNovosCentroidsKernel**). Since the Host implicitly waits for every thread to finish before executing the next Kernel function, this provides the synchronization method we need. It also provides a way to utilize less Threads in the smaller functions, since the resetting and the calculation of the new centroid coordinate phase only require K Threads to run at its full speed.

C. Data Race

On the **calculaSomatorioKernel**, since we have a N number of threads summing/incrementing to a K array, it's inevitable that there will be some Data Races. To control this, we utilized the native CUDA function **atomicAdd(addr,val)**, which ensures that no value can get "lost" amid the sum process, even if it costs some processing time.

D. Architecture



E. Results

CudaV2		
Test	TExec (s)	L1 Cache M.
N=1M, K=320, @SEARCH	1,45	6,51E+6
N=1M, K=1024, @SEARCH	1,71	1,92E+7
N=1M, K=1024, @Laptop	0,73	3,60E+07

We can see we had time reductions in every test performed:

- **N=1M, K=320, @SEARCH:** 1,57s \rightarrow 1,45s (-0,12s)
- **N=1M, K=1024, @SEARCH:** 1,78s \rightarrow 1,71s (-0,07s)
- **N=1M, K=1024, @Laptop:** 0,81s \rightarrow 0,73s (-0,08s)

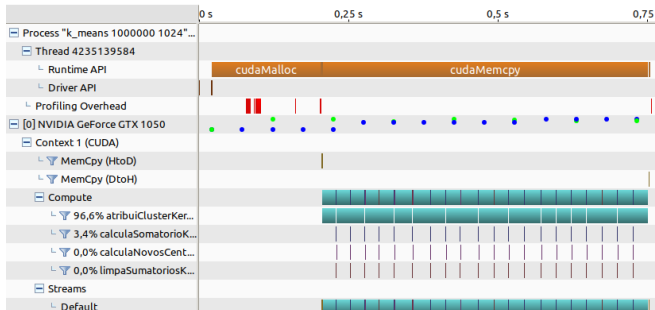
F. Profiling

We proceeded to analyze this version profiling, confirming that this version runs both the algorithm phases:

```

==27874== NVPROF is profiling process 27874, command: /bin/kn_means 1000000 1024
==27874== Profiling application: /bin/kn_means 1000000 1024
==27874== Profiling result:
   Time(s)      Time(s)      Calls      Avg      Min      Max      Name
GPU activities: 96.32s  517.96ns      21  24.665ms  19.756ms  28.855ms  atlbClusterKernel(float)
                3.41s  18.394ms      20  919.69us  741.60us  1.7847ms  calcuSanoClusterKernel(float)
                0.24s  1.2962ms      4   324.04us  1.2160us  653.09us  [CUDA memcpy HtoD]
                0.16s  874.46us      3   291.49us  1.9840us  876.21us  [CUDA memcpy DtoH]
                0.03s  182.34us      20  9.1160us  7.6880us  10.432us  calcuNovosCentroids(float)
                0.03s  146.62ms      20  7.3310ms  6.2680us  8.3700ms  l1pnaSumatorCentroids(int,
API calls:      73.18s  599.13ms      7   77.018ms  9.1170ms  536.33ms  cudaMemCopy
                26.62s  136.13ms      8  24.514ms  2.4130us  195.72ms  cudaMalloc
                0.09s  644.07us      5   128.81us  8.6450us  251.37us  cudaFree
                0.06s  434.15us      81  5.3590us  4.8420us  23.944us  cudaLaunchKernel
                0.03s  326.99us      97  3.3400us  1.6073us  10.377us  cudaGetAttribute
                0.02s  160.37us      1   160.37us  160.37us  160.37us  cudaDeviceTotalMem
                0.01s  45.824us      4   45.824us  45.824us  45.824us  cudaGetDeviceName
                0.00s  8.2160us      1   8.2160us  8.2160us  8.2160us  cudaDeviceGetPCIBusId
                0.00s  2.0040us      3   668ns    163ns    1.0140us  cudaDeviceGetCount
                0.00s  1.9030us      4   490ns    316ns    571ns    cudaDeviceGetError
                0.00s  1.0050us      2   502ns    185ns    820ns    cudaDeviceGet
                0.00s  259ns       1   259ns    259ns    259ns    cudaDeviceGetUuid

```



- The **cudaMalloc** API call had a small execution time increase of around 12ms, due to the increased allocated data in the device for the auxiliary arrays used in the **ajustaCentroids** function.
- The three Kernel functions of **ajustaCentroids**, take up to a combined 19ms, which is around a **100ms reduction** compared to the **123ms CUDA** Downtime present in the last version due to the host running the **ajustaCentroids** function. This reduction is also observed generally in the test result above.
- The number of **cudaMemcpy** API calls is reduced to 7 calls compared to the last version 65 calls. These 7 calls are divided by 4 at the beginning and 3 at the end, only taking up to 2ms of execution time, a **14ms time reduction**.

VI. THIRD VERSION DEVELOPMENT

A. Strategy

With the majority of the algorithm execution running exclusively on the CUDA API, our team decided to look for techniques inside CUDA to raise it’s efficiency and further reduce execution time.

For this, we decided to implement the use of the shared memory present within the CUDA blocks. This memory has a faster read/write access time and is useful for data that is repeatedly accessed/shared between threads.

B. Shared Memory in atribuiClustersKernel

Being the heaviest function even in its CUDA implementation, we started by analyzing how could the use of shared memory improve this function. Since the Point Data and the Attribution Data is generally large in size and its data is accessed in unique positions by every thread, it's an undesirable candidate for this technique.

Meanwhile the centroid data, usually smaller in size, is accessed by all threads, being iterated in its entirety, which makes it a perfect candidate for the use of shared memory. For this, at the beginning of the kernel function, we allocate the shared memory within the block, along with the attribution of the centroid data in this memory. We also have to add a synchronization technique, to ensure the data in the shared memory is all loaded before the rest of the function execution can happen by using the **__syncthreads** CUDA function.

C. Results

CudaV3		
Test	TExec (s)	L1 Cache M.
N=1M, K=320, @SEARCH	1,38	3,92E+6
N=1M, K=1024, @SEARCH	1,50	9,04E+6
N=1M, K=1024, @Laptop	0,48	2,23E+07

In this version, we had a larger performance gain on the tests with a larger number of Clusters. These findings will be discussed in the next chapter.

- **N=1M, K=320, @SEARCH:** 1,45s \rightarrow 1,38s (-0,07s)
- **N=1M, K=1024, @SEARCH:** 1,71s \rightarrow 1,50s (-0,21s)
- **N=1M, K=1024, @Laptop:** 0,73s \rightarrow 0,48s (-0,25s)

D. Profiling


We proceeded to our final profiling, to see how much the kernel functions benefitted from the use of shared memory:

```

==9347== NUPROF is profiling process 9347, command: /bin/kc_means 1000000 1024
==9347== Profiling application: /bin/kc_means 1000000 1024
==9347== Profiling result:

```

Type	Time(s)	Time	Calls	Avg	Min	Max	Name
GPU activities:	99.82%	264.72ms	21	12.581ms	10.949ms	17.145ms	atribu(ClusterKernel(Cl
	4.43%	12.342ms	20	617.08ms	484.06ms	1.8579ms	calcu(scalonKernel(Cl
	0.46%	1.2885ms	4	322.13us	21.690us	646.37us	[CUDA mempty Hot]
	0.27%	746.04us	3	248.68ms	2.1440us	746.86us	[CUDA mempty DtoH]
	0.06%	176.25us	20	8.8120us	1.9940us	1.9940us	calcu(novosCtoshKe
	0.05%	142.27ms	20	7.1130us	6.4600us	7.8800us	lpin(satlonKernel(Cl
	54.74%	370.41ms	39	9.877ms	7.2940us	276.50ms	calcu(scalonKernel(Cl
	44.92%	229.06ms	8	28.632ms	2.5810us	228.66ms	calcu(Malloc
	0.14%	689.01us	5	137.80us	7.6770us	243.32us	cudaFree
	0.09%	440.87us	81	5.4420us	4.6240us	31.311us	cudaLaunchKernel
API calls:	0.07%	359.80ms	97	3.7090ms	188ms	210.12us	cudbveGetAttribute
	0.03%	139.32us	1	139.32us	162.33us	162.33us	cudbveGetAttr
	0.02%	78.540us	78	78.540us	7.2940us	78.540us	cudbveGetName
	0.00%	8.9670ms	1	8.9670ms	8.9670ms	8.9670ms	cudbveGetPCIBusId
	0.00%	8.8590us	45	196ms	149ms	866ms	cudbveGetLastErr
	0.00%	2.2310us	3	743ms	197ms	1.7610ms	cudbveGetCount
	0.00%	1.1060us	2	553ms	271ms	835ms	cudbveGetC
	0.00%	1.1060us	2	553ms	271ms	835ms	cudbveGetC
	0.00%	1.1060us	2	553ms	271ms	835ms	cudbveGetC
	0.00%	1.1060us	2	553ms	271ms	835ms	cudbveGetC



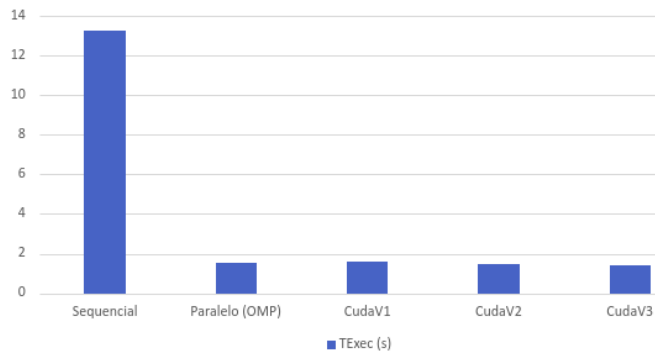
- The **atribuiClusterKernel** function had a 236ms time reduction, representing a close to 50% performance gain on this function.

VII. RESULTS COMPARISON AND ANALYSIS


A. General Comparison

We will start by showing a general comparison of the result on all three tests, including the performance of the sequential and parallelized version developed in previous parts of this project

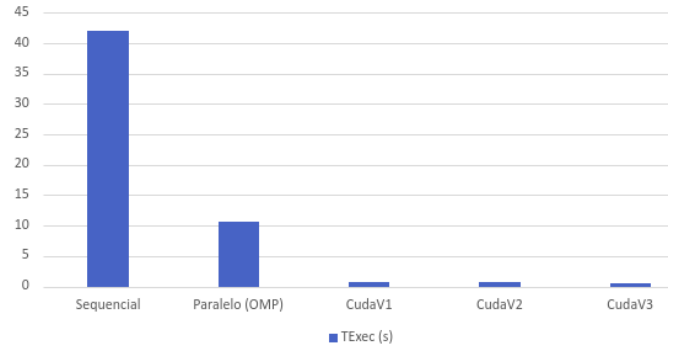
Performance Across Versions
(N=1M, K=320, @ SEARCH)



Performance Across Versions
(N=1M, K=1024, @ SEARCH)



Performance Across Versions
(N=1M, K=1024, @ Laptop)

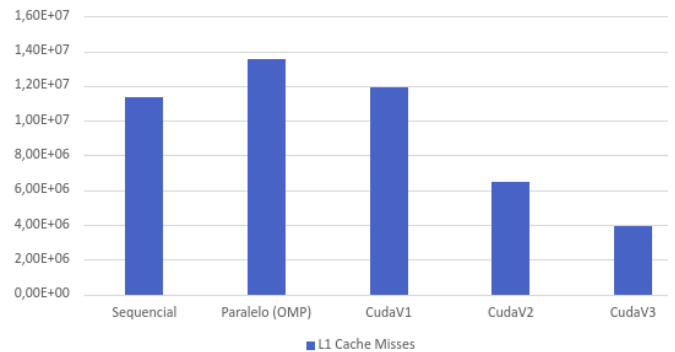


B. L1 Cache Misses

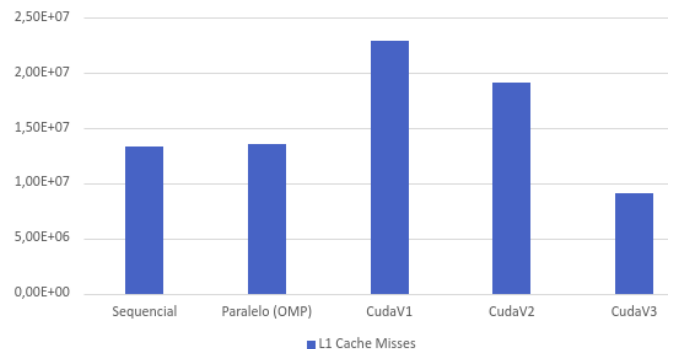
We also analysed the L1 Cache Misses between tests and versions. We can see that within the three CUDA versions developed we can always see a steady decrease in this metric, which reaffirms our observed time reduction.

We can also confirm that the sequential version, although has a lesser level of L1 Cache Misses, has higher execution times since the workload is executed on a single thread, which decreases the likelihood of Cache Misses.

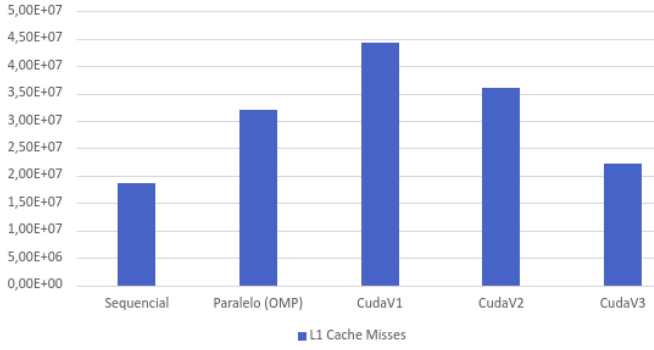
L1 Cache Misses Across Versions
(N=1M, K=320, @ SEARCH)



L1 Cache Misses Across Versions
(N=1M, K=1024, @ SEARCH)



L1 Cache Misses Across Versions
(N=1M, K=1024, @ Laptop)




C. Scalability


To test scalability, we used the two Cluster tests on the SEARCH machine with a K value of 320 and 1024, to analyse how well our CUDA versions scale:

As we can see, compared to the OMP version, which sees a execution time increase of 3s, a 3x increase when the K value changes from 320 to 1024, which is also roughly a 3x increase on the K value, our CUDA version only increase a maximum of 200ms, only a 1,13x execution time increase, which shows the potencial scalability of this solution.

Performance Across Versions
(N=1M, K=320, @ SEARCH)




Performance Across Versions
(N=1M, K=1024, @ SEARCH)



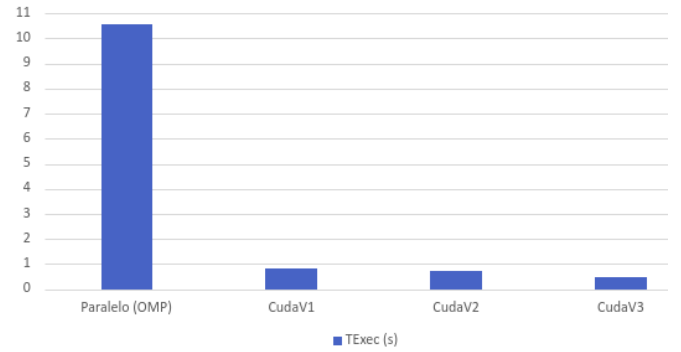
D. SEARCH VS Laptop

In this result set, we can see that the Server machine, with its high core CPU, beats the Laptop on the OMP version, but the Laptop, with its more advanced GPU, has a better performance time than the server on the CUDA versions.

Performance Across Versions
(N=1M, K=1024, @ SEARCH)



Performance Across Versions
(N=1M, K=1024, @ Laptop)



VIII. CONCLUSION

Our team is very satisfied with the results achieved along this project, by improving the algorithm's sequential version, and then improving execution time by using parallelization techniques, including solutions on CPU or GPU, we increased the algorithm performance, observing up to 110x time gain in some of our tests.

IX. ATTACHMENTS

Version	TExec (s)	L1 Cache Misses
Sequencial	13,28	1,14E+07
Paralelo (OMP)	1,53	1,35E+07
CudaV1	1,57	1,19E+07
CudaV2	1,45	6,51E+06
CudaV3	1,38	3,92E+06

$N = 1M, K = 320, @SEARCH$

Version	TExec (s)	L1 Cache Misses
Sequencial	41,69	1,33E+07
Paralelo (OMP)	4,79	1,36E+07
CudaV1	1,78	2,29E+07
CudaV2	1,71	1,91E+07
CudaV3	1,50	9,04E+06

$N = 1M, K = 1024, @SEARCH$

Version	TExec (s)	L1 Cache Misses
Sequencial	42,00	1,87E+07
Paralelo (OMP)	10,56	3,21E+07
CudaV1	0,81	4,42E+07
CudaV2	0,73	3,60E+07
CudaV3	0,48	2,23E+07

$N = 1M, K = 1024, @Laptop$